Introduction to Machine Learning

DSCI6003

galvanize

- What is Machine Learning?
  - What?
  - Why?
  - ML vs. Statistics
- Types of Learning
  - Supervised
  - Unsupervised
  - Semi-Supervised (ML2)
  - Reinforcement
- Supervised Learning
  - 3 Components of a model
  - ML Process

- Lab: Predicting Interest Rates

Field of study that gives computers the ability to learn without being explicitly programmed.

-Arthur Samuel circa 1959

> A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P, if its performance at tasks in T, as measured by P, improves with experience E.
>
> -Tom M. Mitchell

Machine learning is *NOT*:

- Hard coded logic by programmer: **if**s and **else**s...
- Predefined results: completely deterministic
- Burden is placed on programmer at design time
- Must anticipate all inputs to program, and react

Machine learning *is*:

- Automated knowledge acquisition through input
- Iterative improvement as more data is seen
- Adaptive Algorithms

Regression:

- Loan interest rate prediction
- Utilities: smart grid load forecasting
- Web: page traffic prediction
- Advertising CTR prediction

Classification:

- Spam Filtering and document classification
- Finance: Fraud detection and loan default prediction
- Sentiment Analysis: People like to do this with Tweets
- National Security: ??? PRISM!

Clustering:

- Product Marketing: Cohort Analysis
- Oncology: Malignant cell identification
- Computer Vision: entity recognition
- Census: demographics analysis

https://www.kaggle.com/wiki/DataScienceUseCases

Quiz: Which category would Churn prediction fall into?

- Red: Clustering
- Blue: Classification
- Green: Regression
- Yellow: Other

MACHINE LEARNING VS. STATISTICS

Machine learning	Statistics
network, graphs	model
weights	parameters
learning	fitting
generalization	test set performance
supervised learning	regression/classification
unsupervised learning	density estimation, clustering
large grant = $1,000,000	large grant= $50,000
nice place to have a meeting: Snowbird, Utah, French Alps	nice place to have a meeting: Las Vegas in August

http://datavu.blogspot.com/2014/08/statistical-modeling-vs-machine-learning.html

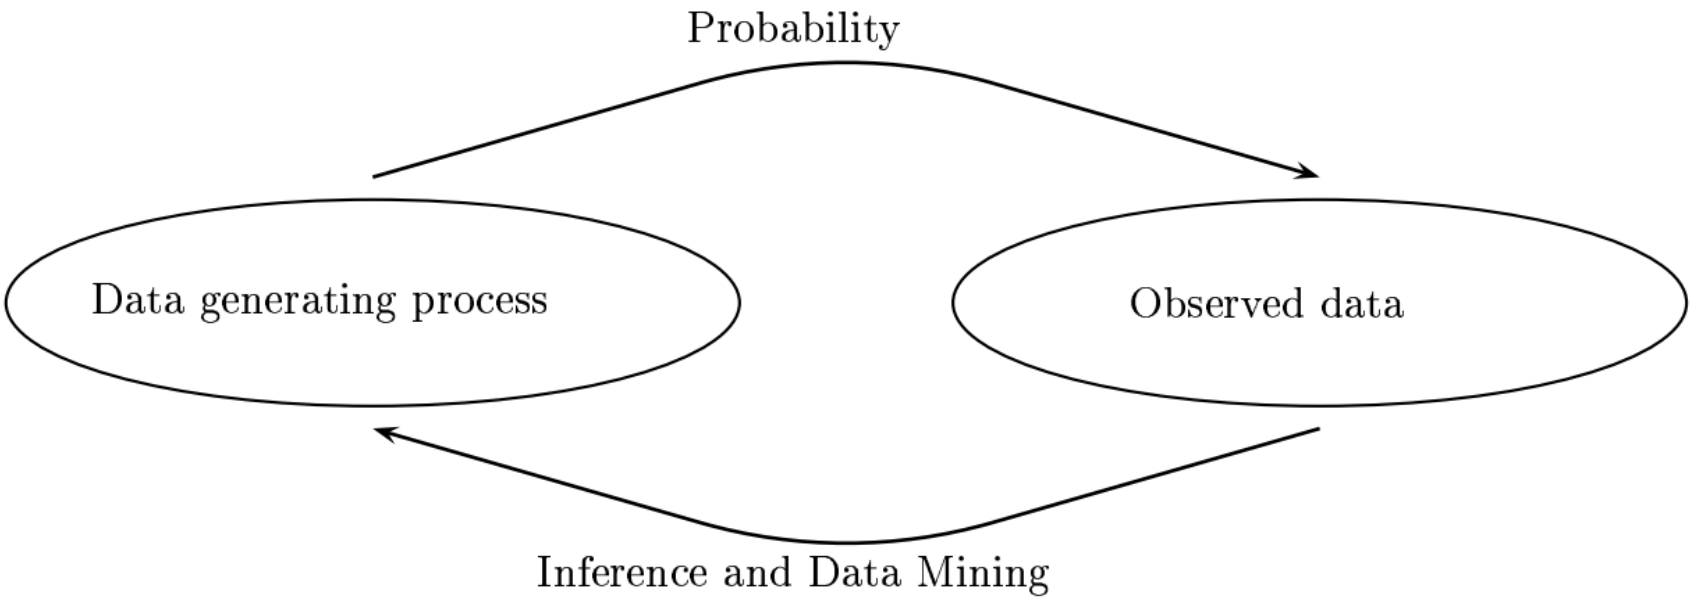
Probability
Data generating process
Observed data
Inference and Data Mining

The Spectrum of the Learning Arts

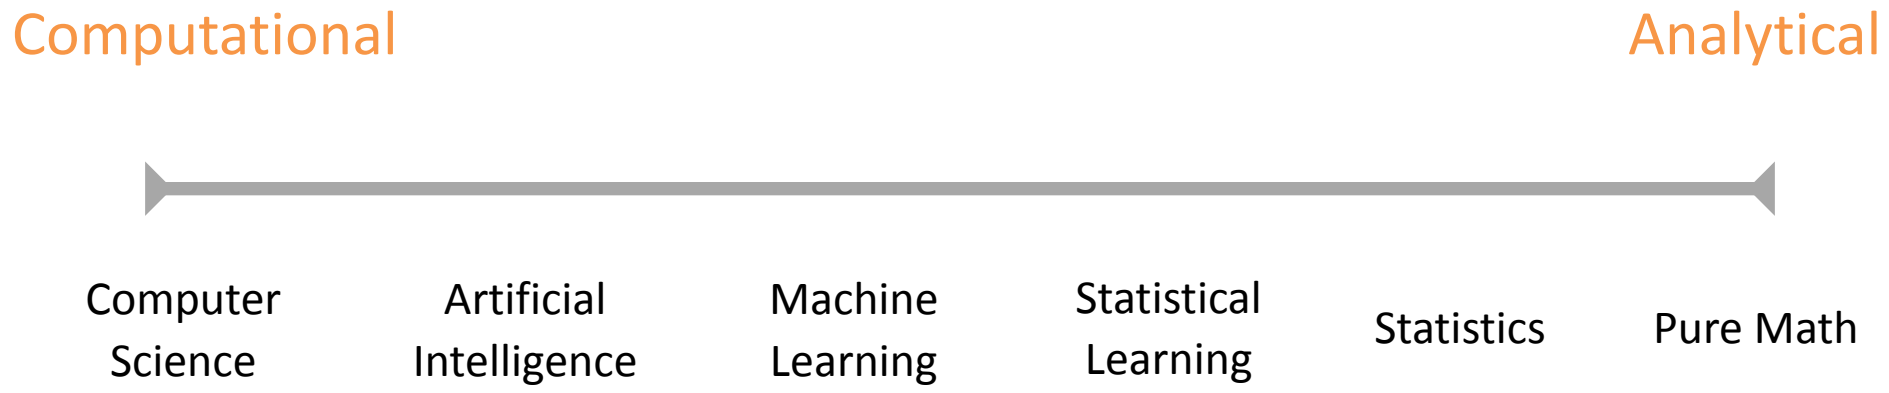

TYPES OF LEARNING

How can computers learn??!?

Supervised Learning

- Training Data includes desired output

Unsupervised Learning

- Training Data does not include desired output

Semi-supervised Learning

- Training Data includes a few desired outputs

Reinforcement Learning

- Rewards from sequence of actions

Supervised Learning

- Training Data includes desired output

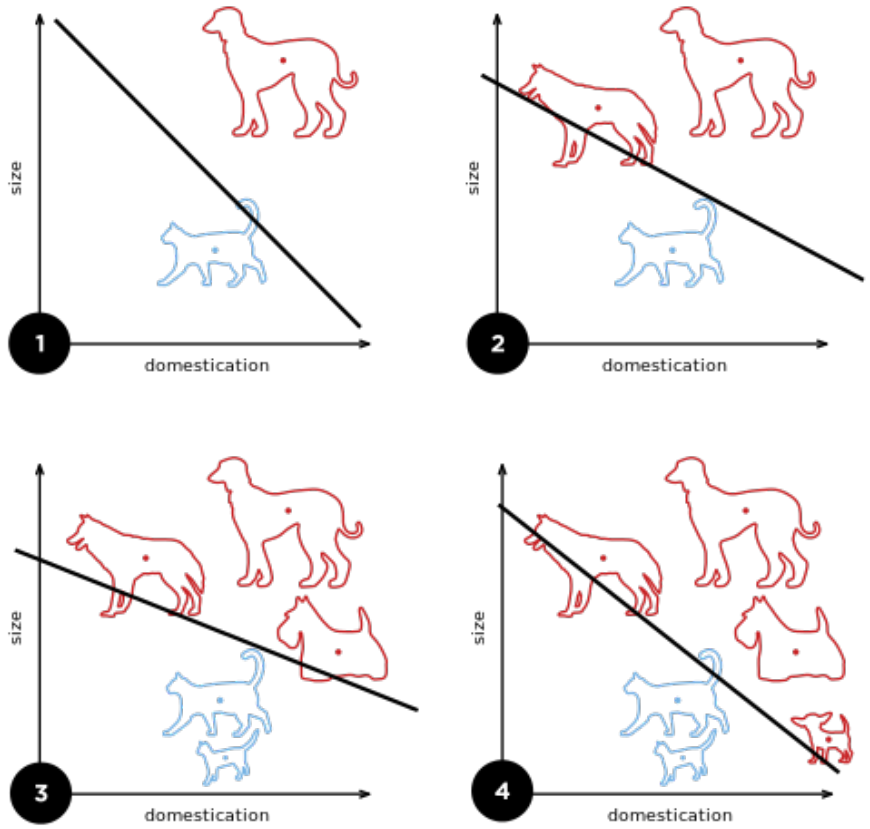

Example: Spam email classifier

Unsupervised Learning

- Training Data does not include desired output

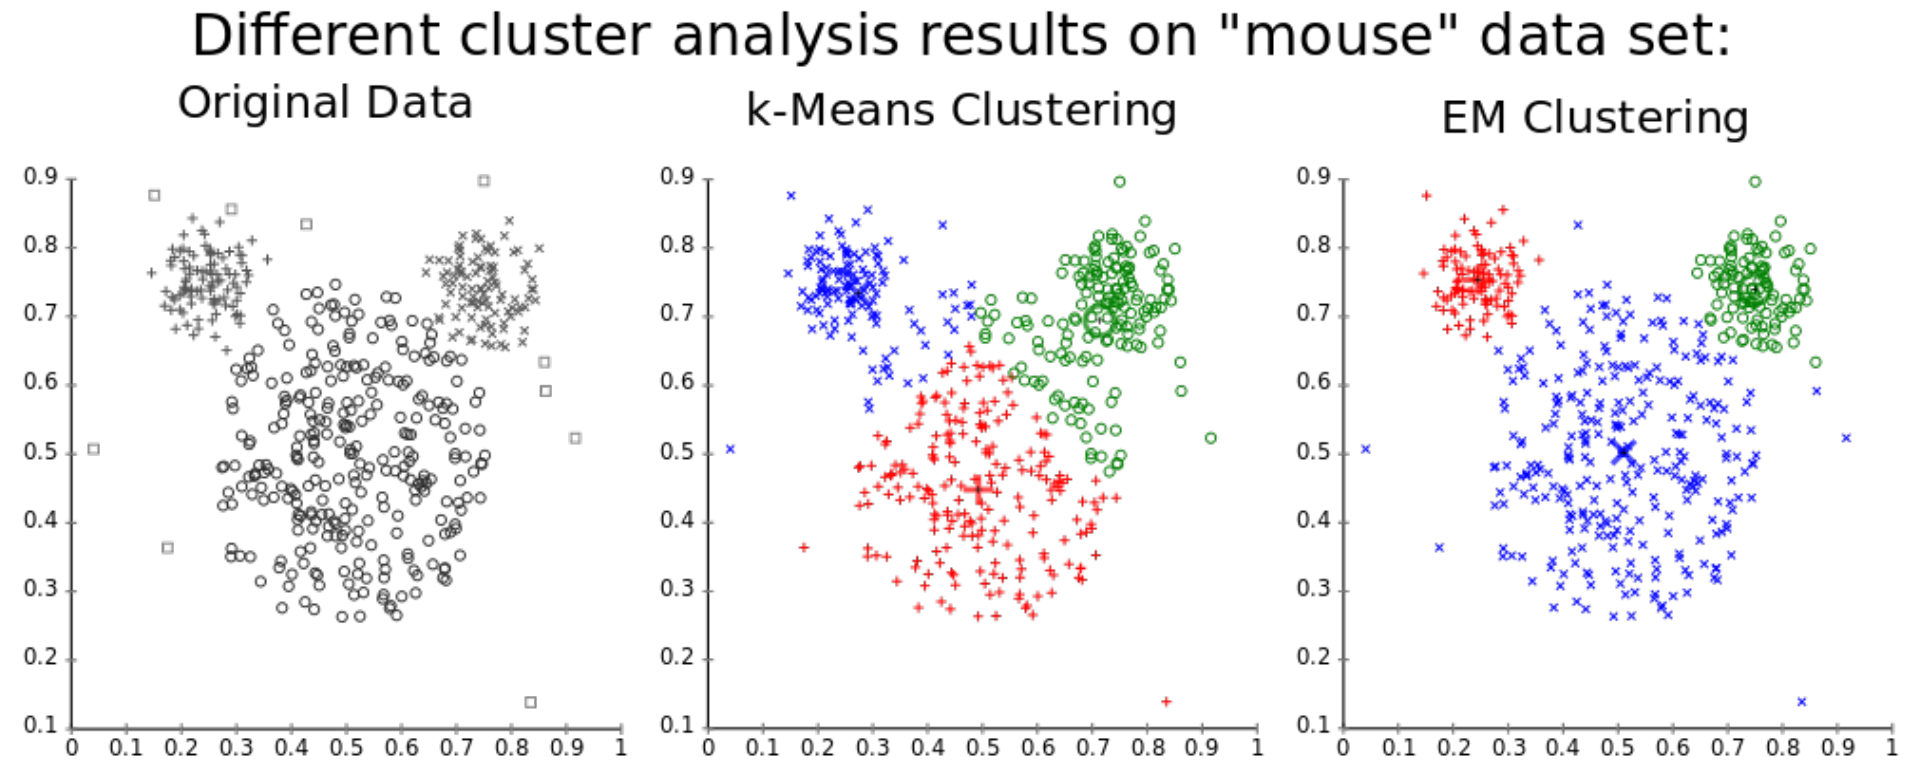

Example: Group users into cohorts based on habits

Semi-supervised Learning

- Training Data includes a few desired outputs

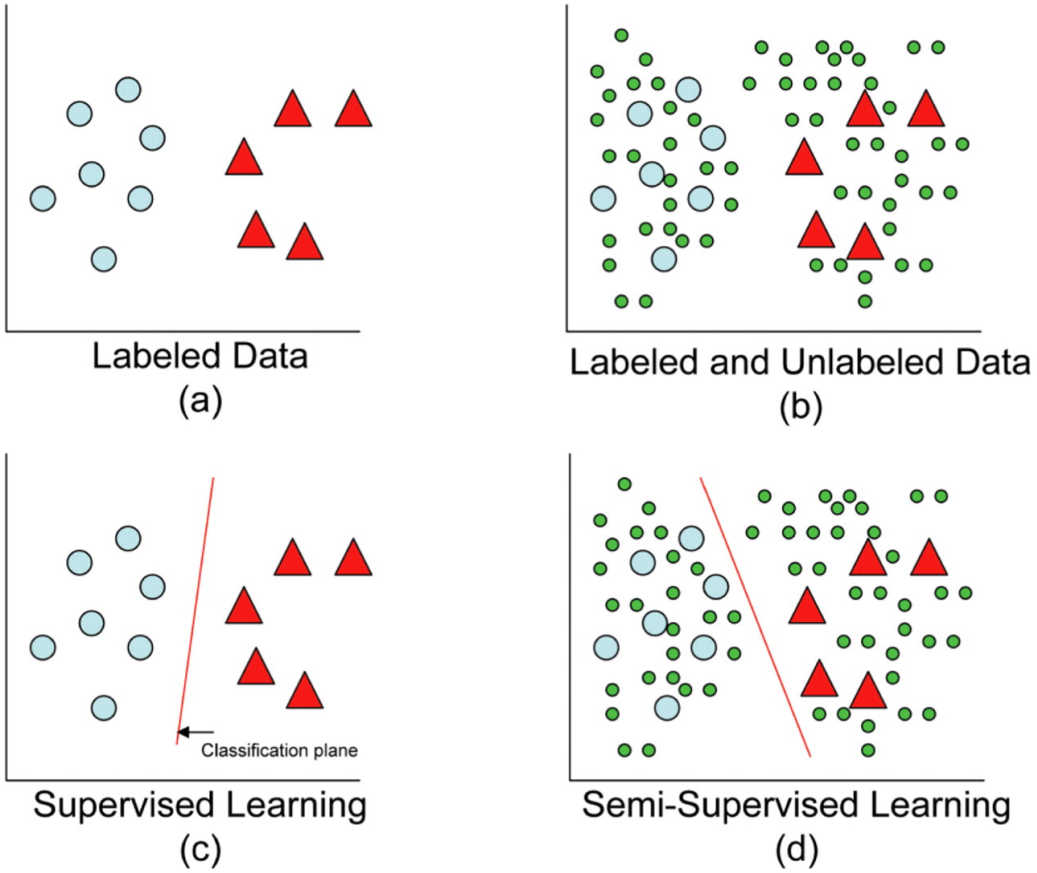

Example: What to do when you don't having enough labels

Reinforcement Learning (traditional AI)

- Rewards from sequence of actions

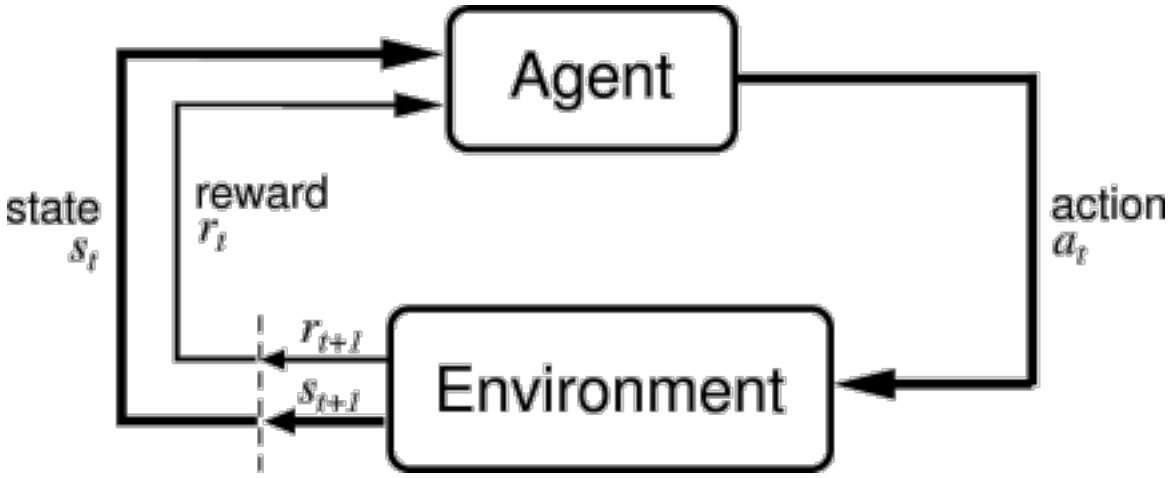

Example: Autonomous Video game Player

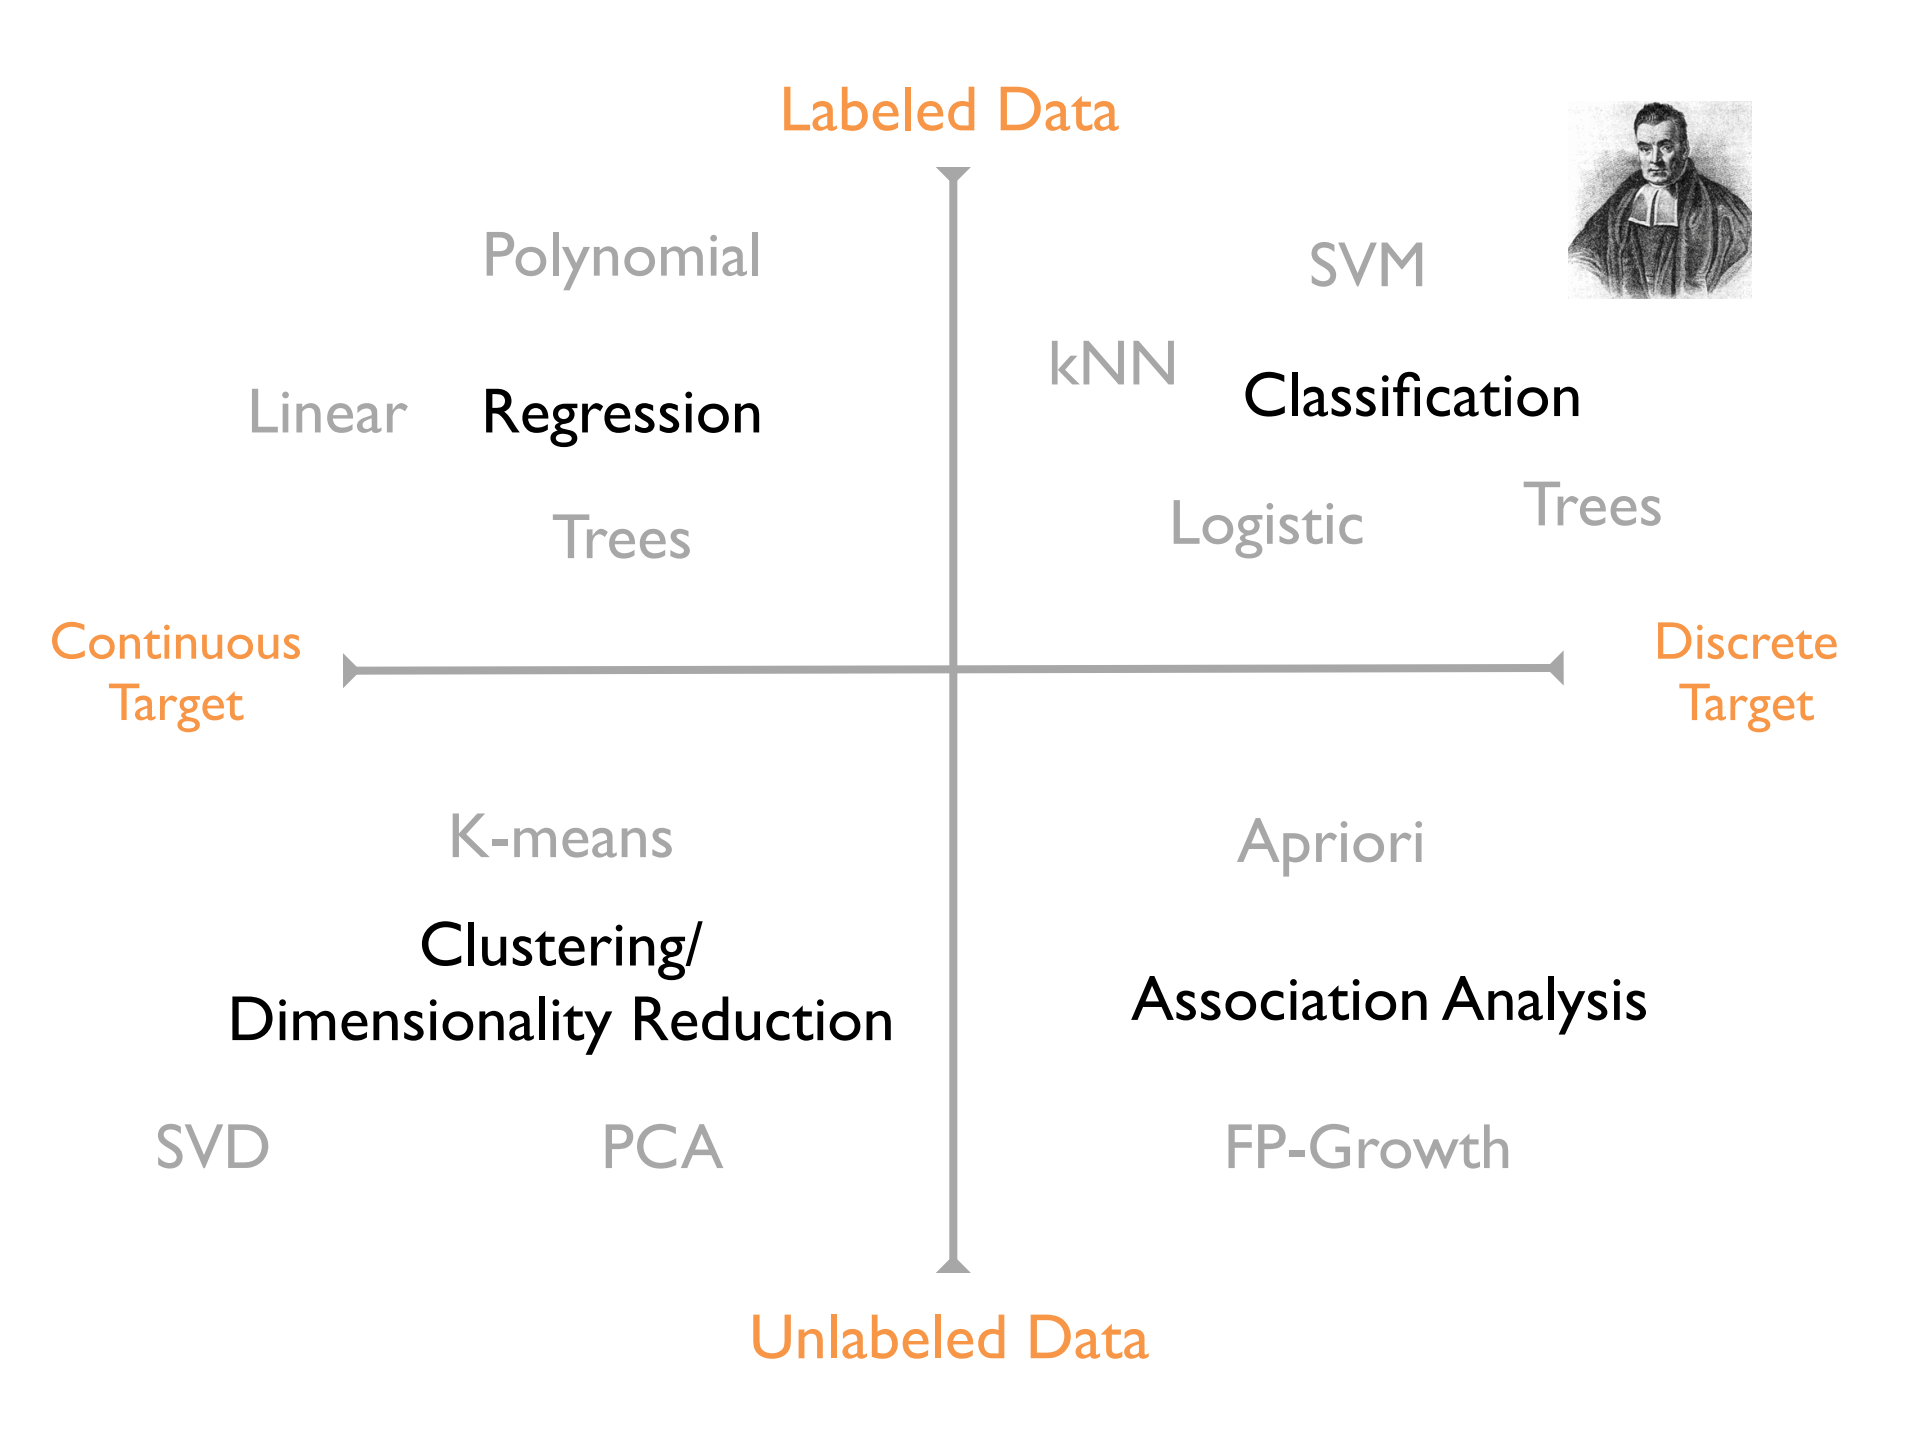
Labeled Data
Polynomial
SVM
kNN
Linear
Regression
Classification
Trees
Logistic
Trees
Continuous Target
Discrete Target
K-means
Apriori
Clustering/ Dimensionality Reduction
Association Analysis
SVD
PCA
FP-Growth
Unlabeled Data

The Unreasonable Effectiveness of Data

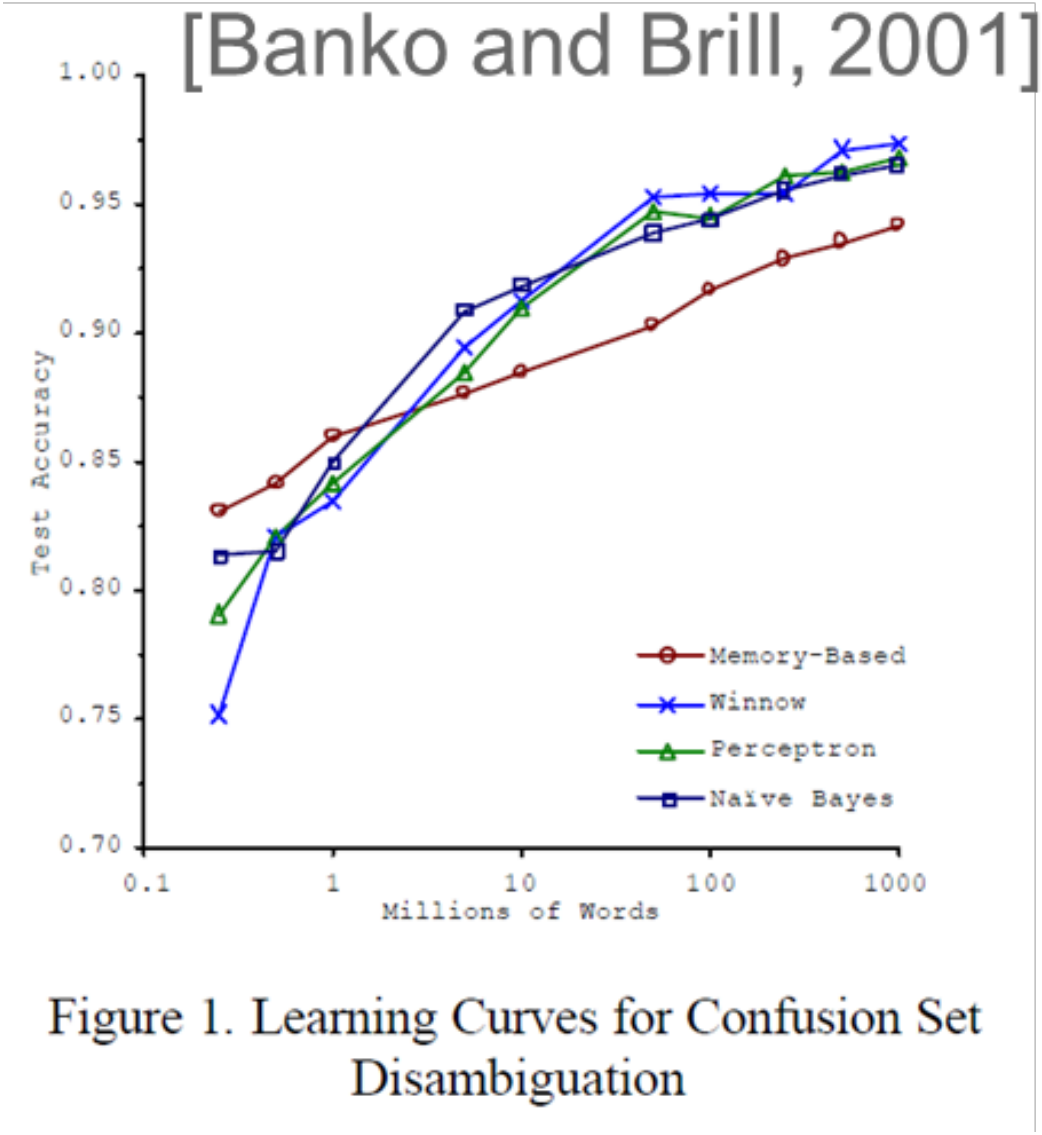

Figure 1. Learning Curves for Confusion Set Disambiguation

Source: http://recsys.acm.org/more-data-or-better-models/

Quiz:

You are the Dean of a college and you need an automated way to send dissertations to the right departments…

Quiz: Which model will you use to avoid having to read all these theses?

- Red: Naive Bayes
- Blue: Linear Regression
- Green: Logistic Regression
- Yellow: Other (neural nets)

Quiz 2:

You are a realtor trying to find undervalued houses on the market…

Quiz 2: Which model will you use to predict the market value of houses?

- Red: Naive Bayes
- Blue: Linear Regression
- Green: Logistic Regression
- Yellow: Other (neural nets)

INTRODUCTION TO SUPERVISED LEARNING

- Supervised Learning
  - 3 Components of a model
    - Hypothesis Function
    - Cost Function
    - Optimization Technique
  - Process
    - Labels vs. Features
    - Train
    - Test
    - Predict
- Lab: Predicting Interest Rates

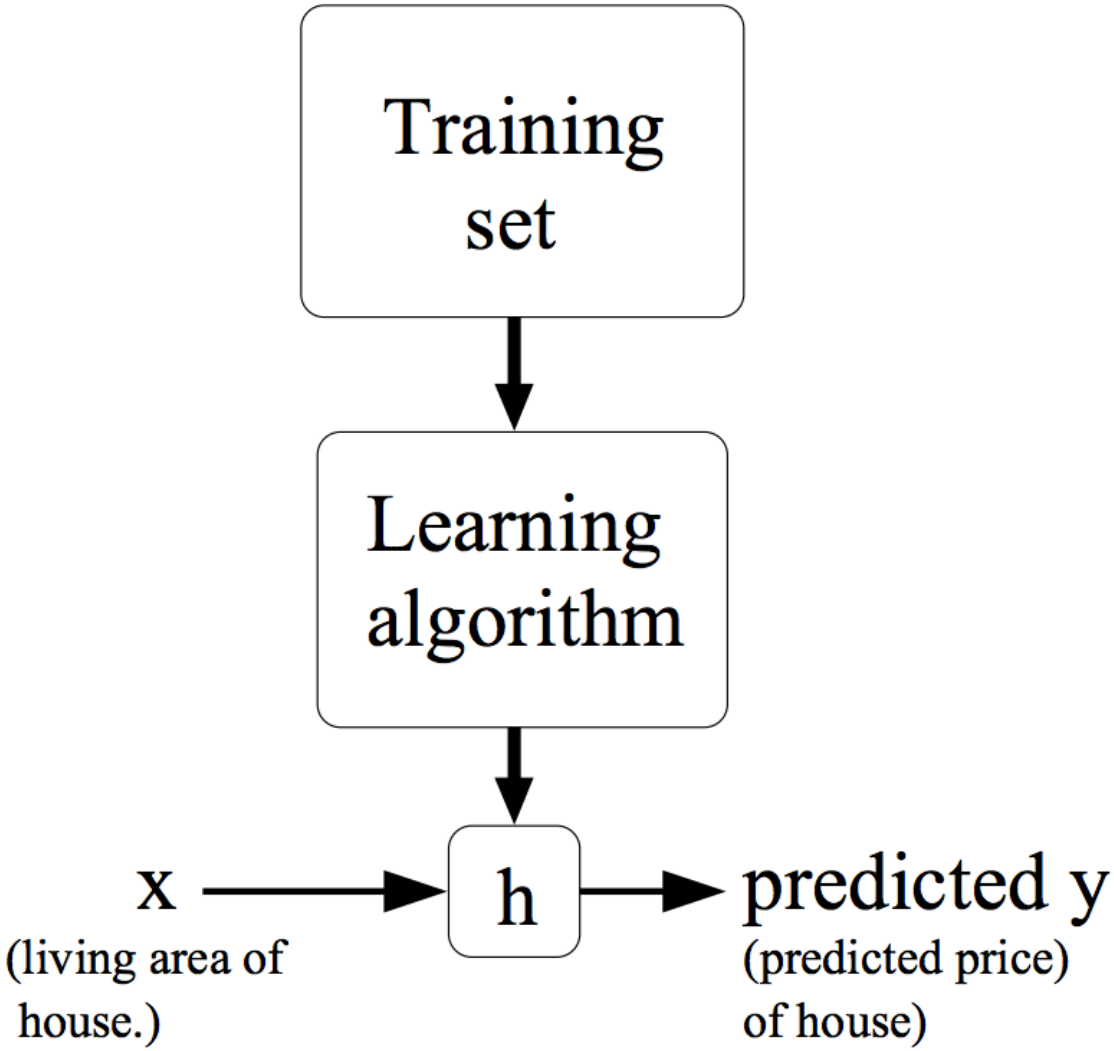

source: Andrew Ng, CS229

Iris Dataset

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	label
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0

Features
(feature matrix)

Target

Train

Input: historical labeled data

+

(hypothesis) function with unknown parameter values
<linear, logistic, etc.>

=

Predict

Output: parameter values

Multi Dimensional Regression

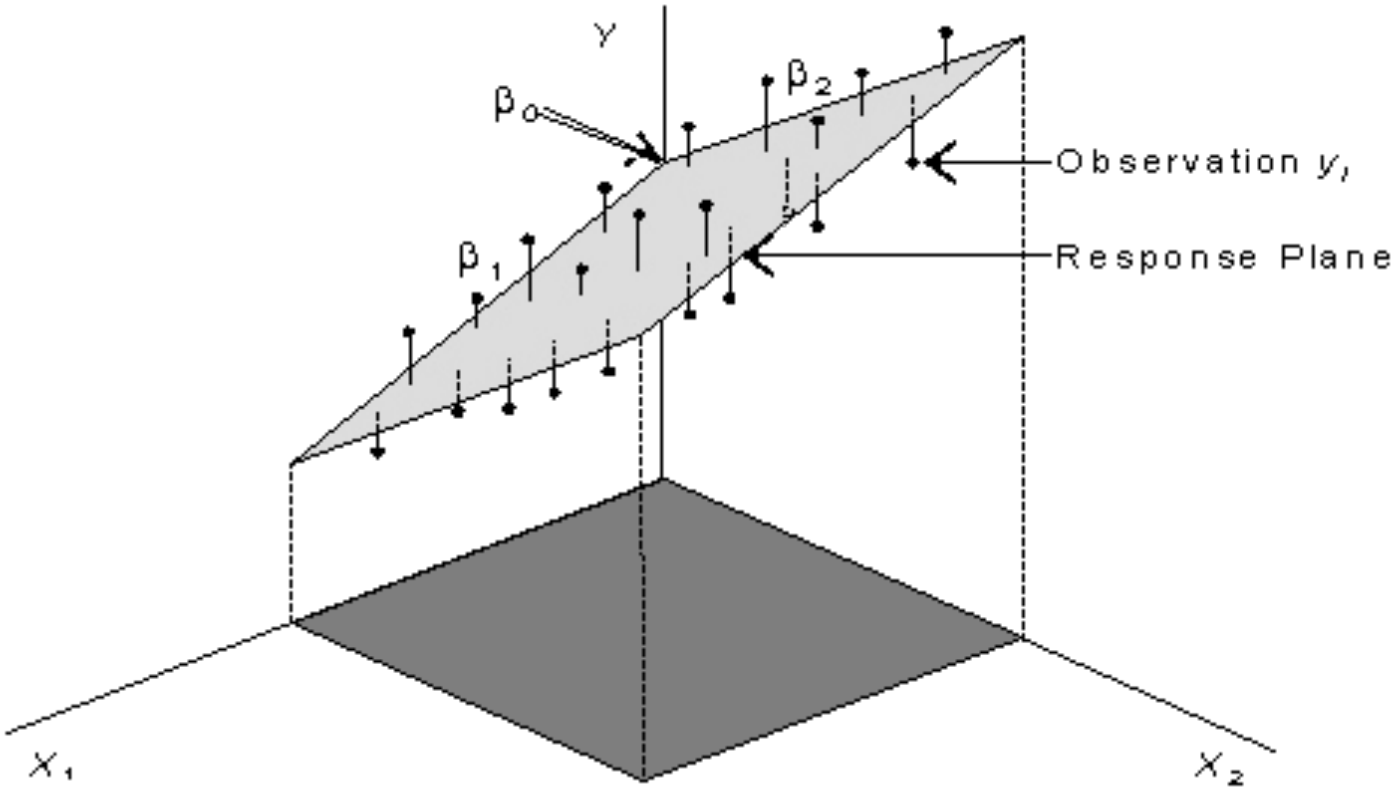

$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \ldots + \beta_n x_n$$

Hypothesis Function

Multi Dimensional Regression

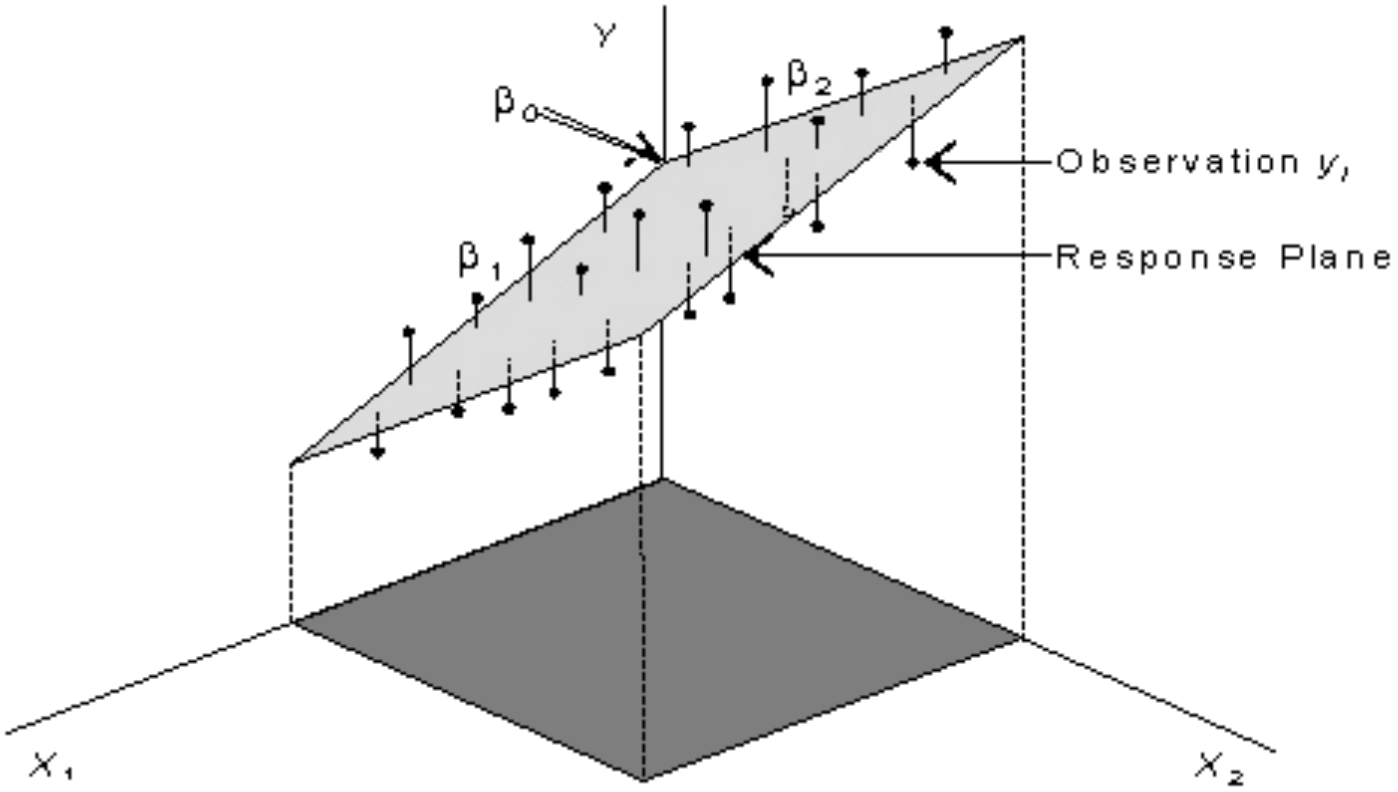

Parameters

$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \ldots + \beta_n x_n$$

Multi Dimensional Regression

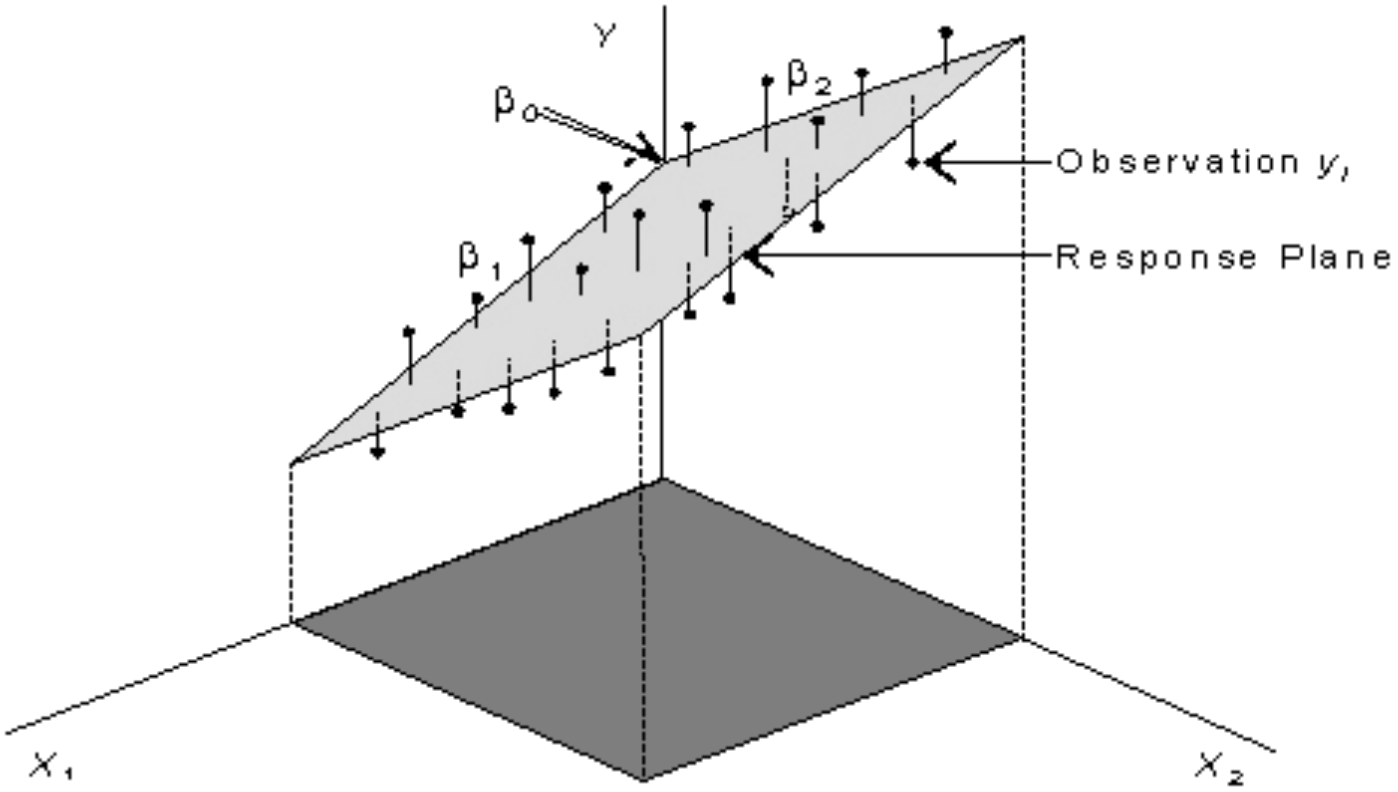

ML Gold

$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \ldots + \beta_n x_n$$

Multi Dimensional Regression

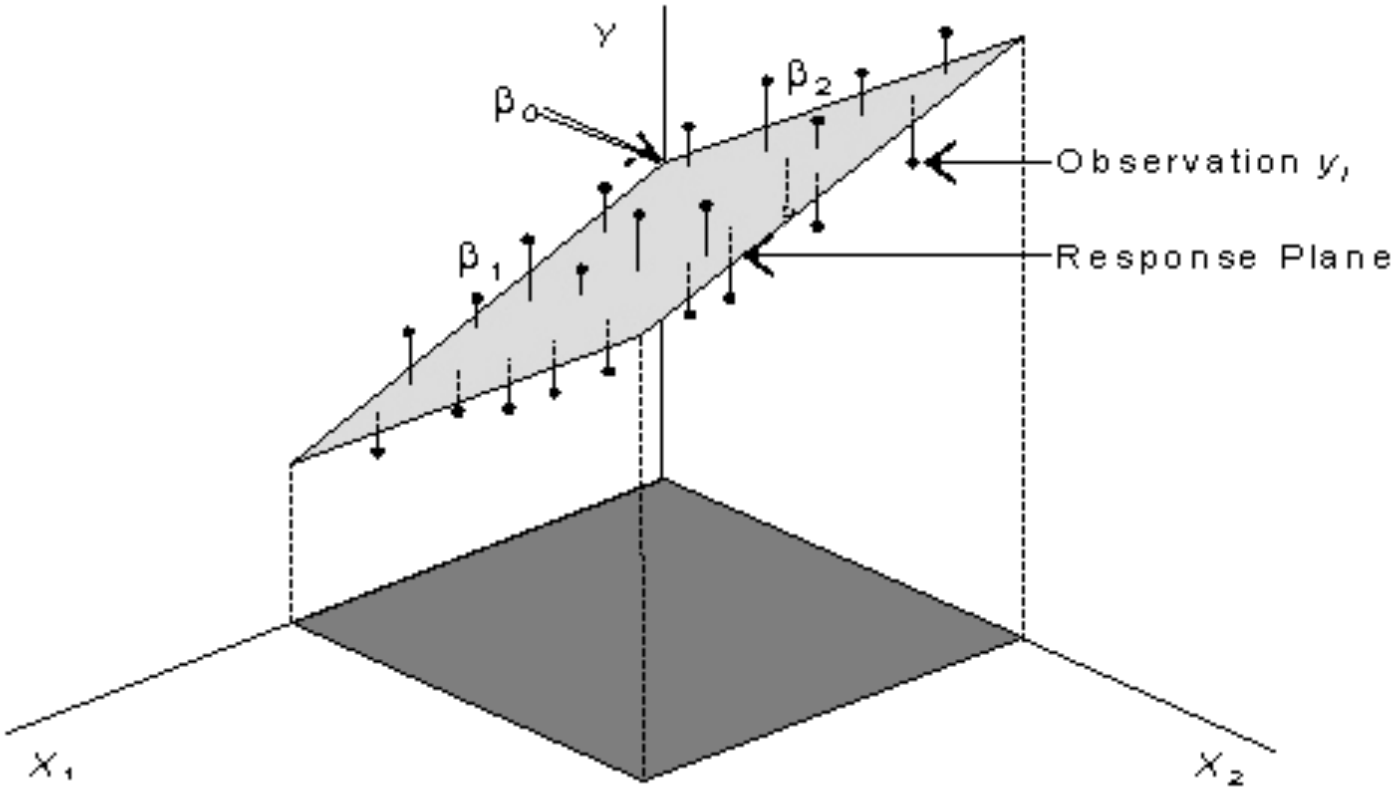

Mine it!

$$\hat{y} = 0.32 + 1.45x_1 + 3.7x_2 + \ldots - 8.91x_n$$

What to learn an unknown target function f()

Input: labeled training set (xi, yi)

- yi = f(xi)

Output: hypothesis h() function “close” to f()

Many possible hypothesis families:

- Logistic
- Linear
- decision trees
- example-based (nearest neighbor)
- etc.

Throughout this class we will work to understand the answers the following questions:

- Which hypothesis space to choose?
- How do we measure goodness of fit?
- How do we balance goodness of fit with complexity?
- How do we make h() a good method?
- How do we pick the right kind of h()?
- How do we know if a good h() will predict well?

Throughout this class we will work to understand the answers the following questions:

- Which hypothesis space to choose?
  - Hypothesis Function
- How do we measure goodness of fit?
  - Cost Function
- How do we balance goodness of fit with complexity?
  - Regularization
- How do we make h() a good method?
  - Optimization
- How do we pick the right kind of h()?
  - Cross Validation
- How do we know if a good h() will predict well?
  - Hold-out Evaluation

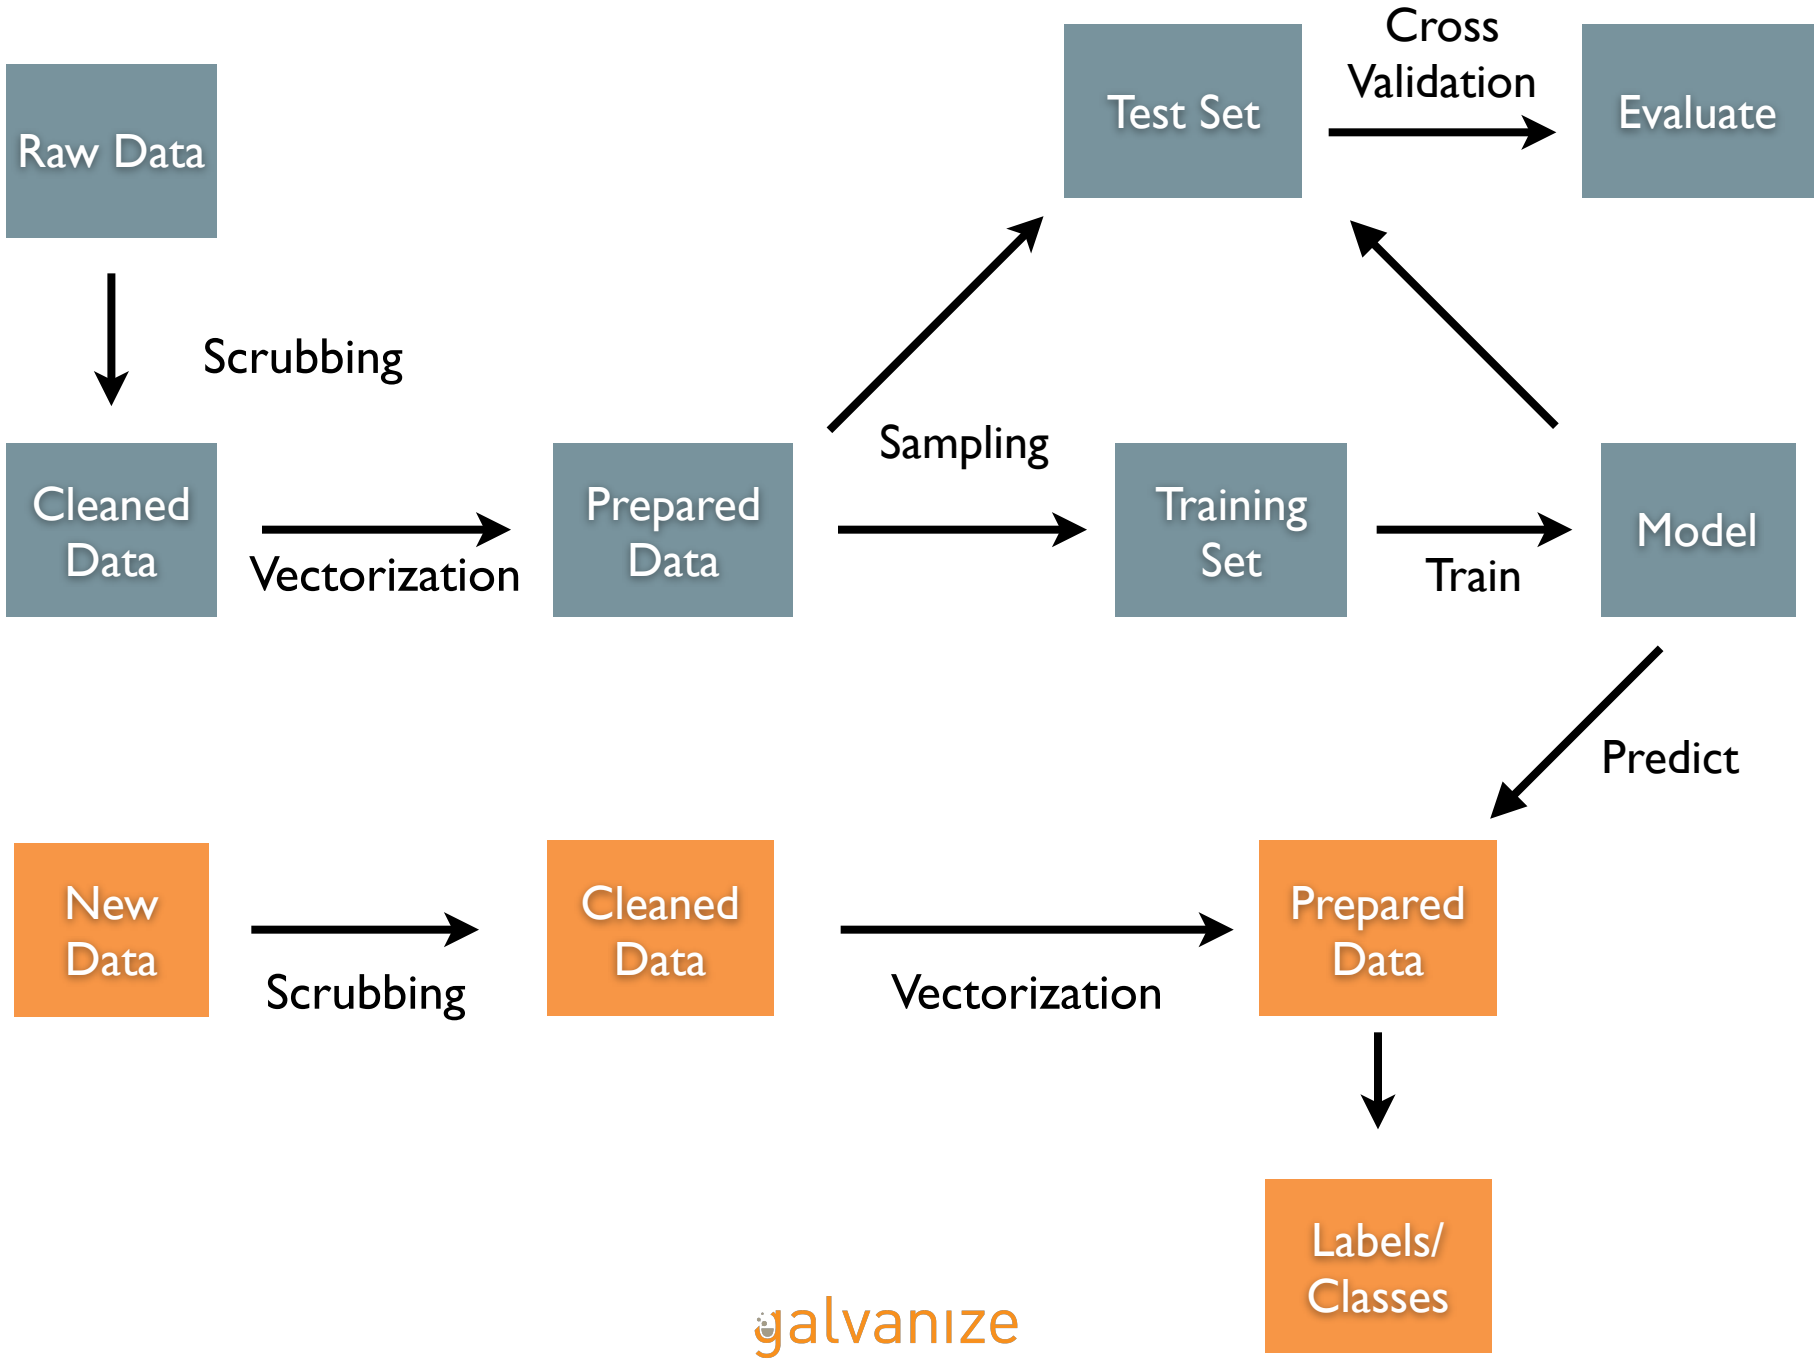
Raw Data
Scrubbing
Cleaned Data
Vectorization
Prepared Data
Sampling
Test Set
Cross Validation
Evaluate
Training Set
Train
Model
Predict
New Data
Scrubbing
Cleaned Data
Vectorization
Prepared Data
Labels/ Classes
galvanize

LAB: PREDICTING LOAN RATES (WITH SCIKIT-LEARN)