

# Bayesian A/B Testing

Beta Distribution and Multi-Arm Bandit

# A/B Testing

(frequentist)

- Define a metric (CTR, for example)
- Determine parameters of interest for study (number of observations, power, significance threshold, and so on)
- Run test, without checking results, until number of observations has been achieved
- Calculate p-value associated with your hypothesis test
- report p-value and suggestion for action

# A/B Testing

(frequentist)

- Can you say “it is 95% likely that site A is better than site B”?
- Can you stop test early based on surprising data?
- Can you update the parameters of your test while it is running?

# A/B Testing

(bayesian)

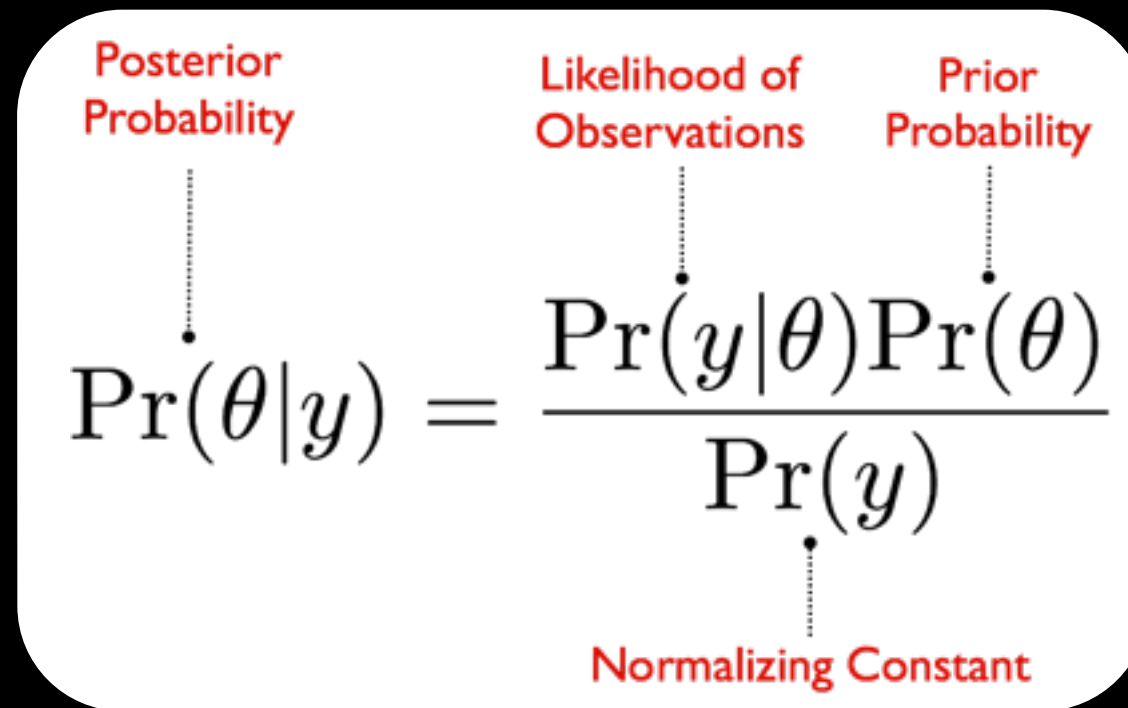
- Define a metric (CTR, for example)
- Run test, continually monitor results
- At any time calculate a probability that site A has better results on the defined metric than site B
- Suggest courses of action based on this probability

# A/B Testing

(bayesian)

- Can you say “it is 95% likely that site A is better than site B”?
- Can you stop test early based on surprising data?
- Can you update the parameters of your test while it is running?

# Bayes Theorem



The diagram shows the Bayes Theorem equation with labels in red text and dotted lines pointing to the corresponding parts of the formula:

- Posterior Probability** points to  $\Pr(\theta|y)$
- Likelihood of Observations** points to  $\Pr(y|\theta)$
- Prior Probability** points to  $\Pr(\theta)$
- Normalizing Constant** points to  $\Pr(y)$

$$\Pr(\theta|y) = \frac{\Pr(y|\theta)\Pr(\theta)}{\Pr(y)}$$

- **prior:** initial belief
- **likelihood:** likelihood of data given outcome
- **posterior:** updated belief

# Bayes Theorem

$$\text{posterior} \propto \text{prior} \times \text{likelihood}$$



# Beta Distribution

$$\frac{p^{\alpha-1} (1-p)^{\beta-1}}{B(\alpha, \beta)}$$

- $p$ : conversion rate (between 0 and 1)
- $\alpha, \beta$ : shape parameters
  - $\alpha = 1 + \text{number of conversions}$
  - $\beta = 1 + \text{number of non conversions}$
- Beta Function (B) is a normalizing constant
- $\alpha = \beta = 1$  gives the *uniform distribution*



# The Distribution



1 = conversion

0 = non conversion

# Binomial (Likelihood)

$$\binom{n}{k} p^k (1 - p)^{n-k}$$

- $p$ : conversion rate (between 0 and 1)
- $n$ : number of visitors
- $k$ : number of conversions

# Conjugate Priors

posterior  $\propto$  prior  $\times$  likelihood

beta  $\propto$  beta  $\times$  binomial

THE MATH:


$$\begin{aligned}\text{posterior} &\propto \text{prior} \times \text{likelihood} \\ &= \frac{p^{\alpha-1}(1-p)^{\beta-1}}{B(a, b)} \times \binom{n}{k} p^k (1-p)^{n-k} \\ &\propto p^{\alpha-1}(1-p)^{\beta-1} \times p^k (1-p)^{n-k} \\ &\propto p^{\alpha+k-1} (1-p)^{\beta+n-k-1}\end{aligned}$$

The result is a Beta Distribution with these shape parameters:

$$\alpha + k \text{ and } \beta + n - k$$

# A/B Testing

- We want to know if this is true:  
**conversion rate of site A > conversion rate of site B**
- We can also answer if this is true:  
**conversion rate of site A > conversion rate of site B + 5%**



Method:

- Sample a large number from both distributions
- Count the percent of times site A wins

# The code

```
num_samples = 10000

A = np.random.beta(1 + num_clicks_A,
                    1 + num_views_A - num_clicks_A,
                    size=num_samples)

B = np.random.beta(1 + num_clicks_B,
                    1 + num_views_B - num_clicks_B,
                    size=num_samples)

### The probability that A wins:
print np.sum(A > B) / float(num_samples)

### The probability that A > B + 0.5%:
print np.sum(A > (B + 0.05)) / float(num_samples)
```

# Multi-Arm Bandit


Smarter A/B Testing

# Traditional A/B Testing

- First: pure *exploration*, in which you assign equal numbers of users to Group A and Group B
- Second: pure *exploitation*, in which you stop the experiment and send all your users to the more successful version of your site.

# Multi-Arm Bandit

- Start *exploiting* the likely best solution *before* you're done *exploring*
- versions of the site  
= **bandits** (slot machines)
- How to pick which one to play....





# Regret

- To evaluate a multi-arm bandit algorithm, *minimize regret*
- **Regret** is a measure of how often you chose a suboptimal bandit

$$\begin{aligned}\text{regret} &= \sum_{i=1}^k (p_{\text{opt}} - p_i) \\ &= k \cdot p_{\text{opt}} - \sum_{i=1}^k p_i\end{aligned}$$

# Epsilon-Greedy Algorithm

- Explore with probability epsilon (often 10%)
- Exploit the rest of the time (play the bandit with the best performance so far)

# UCB1 Algorithm (Upper Confidence Bound)

- Choose the bandit where this value is the largest

$$p_A + \sqrt{\frac{2 \log N}{n_A}}$$

where  $n_A$  = number of times bandit A has been played  
and  $N$  = total number of times any bandit has been played

# Softmax Algorithm

- Choose the bandit randomly in proportion to its estimated value:

$$\frac{e^{p_A/\tau}}{e^{p_A/\tau} + e^{p_B/\tau} + e^{p_C/\tau}}$$

# Bayesian Bandit Algorithm

- Model each of the bandits with a beta distribution with shape parameters:  
 $\alpha = 1 + \text{number of times bandit has won}$   
 $\beta = 1 + \text{number of times bandit has lost}$
- Take a random sample from each bandit's distribution and choose the bandit with the highest value.

# Bayesian Bandit Simulation

3 bandits, payouts of  $[0.05, 0.1, 0.2]$

