

# Monte Carlo Integration

Robert Winslow

GalvanizeU

June 20th, 2017

## Warm-up discussion

What does it mean to *describe* a probability distribution?

What does it mean to *use* a probability distribution?

# Learning objectives

- ▶ Know when you should consider approximation methods. (Lecture)
- ▶ Build an intuition for Monte Carlo integration in general. (Lecture + Lab)
- ▶ Understand MC integration for approximating expectations. (Lecture + Lab)
- ▶ Understand MC error. (Lab)
- ▶ Understand how naive MC integration can be impractical. (Lab)

# Model computation

Recall the definition of a probability distribution  $P$ :

- ▶ Non-negativity:  $P(A) \geq 0$
- ▶ Normalization:  $P(\Omega) = 1$
- ▶ Finite additivity:  $A \cap B = \emptyset \rightarrow P(A \cup B) = P(A) + P(B)$

What does the definition of  $P$  say about how to *compute* with  $P$ ?

Nothing.

# Model computation

How to compute with  $P$ , and with random variables in general, is up to us.

Models can be arbitrarily complicated.

Computation techniques will be specific to the implementation of a model.

# Querying

In particular, we care about how to *query* our model.

Recall some query types for a random variable  $X$ :


- ▶ Expectation:  $E[X]$
- ▶ Draw:  $a \sim X$
- ▶ PDF:  $P(X = a)$

Sometimes, querying is easy.

But, in general, querying is difficult.


Computer science note: easy  $\sim \mathcal{O}(1)$ , difficult  $> \mathcal{O}(1)$ , very difficult  $\geq \mathcal{O}(2^n)$ .

## Example: Continuous uniform distribution $\mathcal{U}(a, b)$



- ▶  $E[X]$ : Closed form:  $\frac{a+b}{2}$  (easy)
- ▶  $a \sim X$ : Given by computing environment. (easy)
- ▶  $P(X = a)$ : Closed form:  $\frac{1}{b-a}$  (easy)

## Example: Normal distribution $\mathcal{N}(\mu, \sigma^2)$




- ▶  $E[X]$ : Closed form:  $\mu$  (easy)
- ▶  $a \sim X$ : Analytic: Box-Muller, Ziggurat, etc. (easy)
- ▶  $P(X = a)$ : Closed form:  $\frac{e^{-\frac{(x-\mu)^2}{2\sigma^2}}}{\sqrt{2\pi\sigma^2}}$  (easy)

N.B. Box-Muller, Ziggurat, and related methods are a special case for this distribution.



# Example: Bayesian networks



- ▶  $E[X]$ : Algorithmic: exponential in the treewidth (easy to very difficult)
- ▶  $a \sim X$ : Algorithmic: depends on conditioned variables (easy to very difficult)
- ▶  $P(X = a)$ : Algorithmic: chain rule on the network (medium)

N.B. Bayes models can be easy to very difficult.

# Most distributions are not amenable to querying directly

N.B. Even the humble normal distribution requires a *specially-derived* approach for sampling.

Two big ideas:

- ▶ A few well-studied models have *easy* (closed-form or analytic) query procedures.
- ▶ Querying most models is *difficult* and requires approximations.

Questions so far?



# Our approach

This week, we will examine several approximation techniques.  
We start with *Monte Carlo Integration*.

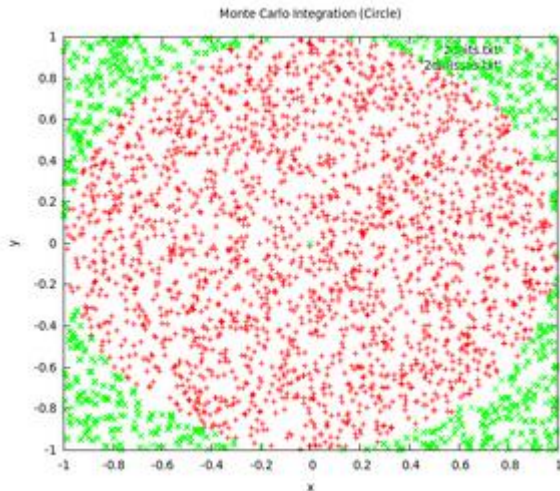
# Monte Carlo Integration

MCI is a general technique to approximate the mean of an integral.

We use it when we can *draw* from  $X$  but we cannot calculate the *expectation* of  $X$ .

You have seen MCI before.

# MCI: area of a circle



area of circle = volume of square \* fraction of points that land in the circle

“Dartboard method”

# MCI: History

The “Monte Carlo” method was developed by Ulam and von Neumann during the Manhattan Project.

Named after Monte Carlo, Monaco. Where there is supposedly a lot of gambling.

Used to create stochastic simulations to build “better” atomic bombs.

# MCI: In general

*MCI gives approximate integrals.*

$I$  is the expression to integrate:

Exact:

$$I = \int_a^b f(x) dx$$

Approximate:

$$\hat{I} = (b - a)MCI(f(x))$$

Monte Carlo Integration can be used when we want to compute the integral, but we can't do it analytically.



# MCI: In statistics

We use MCI to approximate the expectation when exact computation is difficult:

$E[X]$	difficult
$a \sim X$	easy
$P(X = a)$	easy

Exact:

$$E[X] = \int pdf(X)Xdx$$

Approximate:

$$\widehat{E[X]} = MCI(X)$$

# MCI: algorithm

To approximately find the mean of the function  $f(x)$  with MCI:

- ▶ Draw  $N$  samples from  $f$ .
- ▶ Average those samples.

Straightforward!

But, make sure you sample correctly.

## MCI: worked example

$$I = \int_0^{10} x^3 dx$$

```
import numpy as np
import scipy.integrate as integrate
def f(x): return x**3
```

```
given_area = integrate.quad(f, 0, 10)[0]
given_mean = given_area / 10.
```

```
samples = f(np.random.uniform(0, 10, 1000))
approx_mean = np.mean(samples)
approx_area = approx_mean * 10
print(approx_mean, given_mean)
250.785410389 250.00000000000006
```